

Postgres Lifecycle Management Operators in Kubernetes

Miroslav Šiřina

Bachelor's thesis
2023



Tomas Bata University in Zlín
Faculty of Applied Informatics

Univerzita Tomáše Bati ve Zlíně
Fakulta aplikované informatiky
Ústav informatiky a umělé inteligence

Akademický rok: 2022/2023

ZADÁNÍ BAKALÁŘSKÉ PRÁCE

(projektu, uměleckého díla, uměleckého výkonu)

Jméno a příjmení: **Miroslav Šiřina**
Osobní číslo: **A20079**
Studijní program: **B0613A140020 Softwarové inženýrství**
Forma studia: **Kombinovaná**
Téma práce: **Operátoři pro správu životního cyklu databázového systému Postgres v Kubernetes**
Téma práce anglicky: **Postgres Lifecycle Management Operators in Kubernetes**

Zásady pro vypracování

1. Seznamte se s operátory v Kubernetes a Postgres databázovým systémem.
2. Definujte životní cyklus databázového serveru.
3. Vyberte vhodné operátory a popište je.
4. Sestavte metodiku testování jednotlivých operátorů.
5. Na základě sestavené metodiky otestujte vybrané operátory.
6. Proveďte zhodnocení.

Forma zpracování bakalářské práce: **tištěná/elektronická**
Jazyk zpracování: **Angličtina**

Seznam doporučené literatury:

1. SAYFAN, Gigi, 2020. Mastering Kubernetes: level up your container orchestration skills with Kubernetes to build, run, secure, and observe large-scale distributed apps. Third edition. Birmingham: Packt Publishing. ISBN 9781839211256.
2. RIGGS, Simon a Gianni CIOILLI, 2022. PostgreSQL 14 Administration Cookbook. Birmingham: Packt publishing. ISBN 9781803248974.
3. DOBIES, Jason a Joshua WOOD. Kubernetes operators: automating the container orchestration platform. Sebastopol, CA: O'Reilly Media, 2020. ISBN 9781492048046.
4. FARLEY, David, 2021. Modern software engineering: doing what really works to build better software faster. Boston: Addison-Wesley. ISBN 9780137314911.
5. DAME, Michael, 2022. The Kubernetes Operator Framework Book. Birmingham: Packt Publishing. ISBN 9781803232850.

Vedoucí bakalářské práce: **Ing. Peter Janků, Ph.D.**
Ústav informatiky a umělé inteligence

Datum zadání bakalářské práce: **2. prosince 2022**
Termín odevzdání bakalářské práce: **26. května 2023**



doc. Ing. Jiří Vojtěšek, Ph.D. v.r.
děkan

prof. Mgr. Roman Jašek, Ph.D., DBA v.r.
ředitel ústavu

Ve Zlíně dne 7. prosince 2022

I hereby declare that:

- I understand that by submitting my Bachelor's thesis, I agree to the publication of my work according to Law No. 111/1998, Coll., On Universities and on changes and amendments to other acts (e.g. the Universities Act), as amended by subsequent legislation, without regard to the results of the defence of the thesis.
- I understand that my Bachelor's Thesis will be stored electronically in the university information system and be made available for on-site inspection, and that a copy of the Bachelor's Thesis will be stored in the Reference Library of the Faculty of Applied Informatics, Tomas Bata University in Zlín.
- I am aware of the fact that my Bachelor's Thesis is fully covered by Act No. 121/2000 Coll. On Copyright, and Rights Related to Copyright, as amended by some other laws (e.g. the Copyright Act), as amended by subsequent legislation; and especially, by §35, Para. 3.
- I understand that, according to §60, Para. 1 of the Copyright Act, Tomas Bata University in Zlín has the right to conclude licensing agreements relating to the use of scholastic work within the full extent of §12, Para. 4, of the Copyright Act.
- I understand that, according to §60, Para. 2, and Para. 3, of the Copyright Act, I may use my work – Bachelor's Thesis, or grant a license for its use, only if permitted by the licensing agreement concluded between myself and Tomas Bata University in Zlín with a view to the fact that Tomas Bata University in Zlín must be compensated for any reasonable contribution to covering such expenses/costs as invested by them in the creation of the thesis (up until the full actual amount) shall also be a subject of this licensing agreement.
- I understand that, should the elaboration of the Bachelor's Thesis include the use of software provided by Tomas Bata University in Zlín or other such entities strictly for study and research purposes (i.e. only for non-commercial use), the results of my Bachelor's Thesis cannot be used for commercial purposes.
- I understand that, if the output of my Bachelor's Thesis is any software product(s), this/these shall equally be considered as part of the thesis, as well as any source codes, or files from which the project is composed. Not submitting any part of this/these component(s) may be a reason for the non-defence of my thesis.

I herewith declare that:

- I have worked on my thesis alone and duly cited any literature I have used. In the case of the publication of the results of my thesis, I shall be listed as co-author.
- The submitted version of the thesis and its electronic version uploaded to IS/STAG are both identical.

In Zlín; dated:

.....

Student's Signature

ABSTRAKT

Text abstraktu česky

Klíčová slova: Přehled klíčových slov

ABSTRACT

Text of the abstract

Keywords: Some keywords

Zde je místo pro případné poděkování, motto, úryvky knih, básní atp.

TABLE OF CONTENTS

INTRODUCTION.....	8
I THEORY	9
1 BACKGROUND	10
1.1 POSTGRES.....	10
1.2 KUBERNETES	11
1.2.1 Kubernetes components	12
1.2.2 Kubernetes Concepts	13
1.3 RUNNING POSTGRES IN KUBERNETES.....	13
1.4 ARCHITECTURE.....	13
2 NADPISY A PODNADPISY	14
2.1 PODNADPIS A	14
2.2 PODNADPIS B	14
2.3 PODNADPIS C	14
2.3.1 Podpodnadpis alfa	14
2.3.2 Podpodnadpis beta	14
2.3.3 Podpodnadpis gama.....	14
2.4 PODNADPIS D	14
3 VKLÁDÁNÍ OBRÁZKŮ, TABULEK A CITACÍ.....	15
3.1 OBRÁZEK.....	15
3.2 TABULKA	15
II PRAKTICKÁ ČÁST	16
4 NADPIS PRVNÍ KAPITOLY PRAKTICKÉ ČÁSTI.....	17
ZÁVĚR.....	18
REFERENCES	19
LIST OF ABBREVIATIONS.....	20
LIST OF FIGURES.....	21
LIST OF TABLES.....	22
LIST OF APPENDICES	23

INTRODUCTION

TBD - První řádek prvního odstavce v kapitole či podkapitole se neodsazuje, ostatní ano. Vertikální odsazení mezy odstavci je typické pro anglickou sazbu; czech babel toto respektuje, netřeba do textu přidávat jakékoliv explicitní formátování, viz ukázka sazby tohoto textu s následujícím odstavcem).

Formátování druhého odstavce. Text text text text text text text text text text text.

I. THEORY

1 BACKGROUND

This chapter introduces the key technologies used in this thesis including Postgres, Kubernetes, and Kubernetes operators.

1.1 Postgres

PostgreSQL is a powerful object-relational database management system (ORDBMS) derived from the POSTGRES package written at the University of California at Berkeley. [1] [2] The first version of POSTGRES was released in June 1989. POSTGRES has been used in many applications, including financial data analysis systems, asteroid tracking databases, medical information database, and several geographic information systems. The size of external community users has nearly doubled by 1993. [3]

POSTGRES was using its POSTQUEL query language from version one until 1995, when Andrew Yu and Jolly Chen introduced SQL to POSTGRES. The name has changed to Postgres95. Postgres95 was completely ANSI C code reduced by 25 % and was 30 – 50 % faster than Postgres 4.2. [3]

It was clear by 1996 that the name would not stand the test of time therefore it has been renamed to PostgreSQL. As stated by PostgreSQL documentation [3]: “Many people continue to refer to PostgreSQL as “Postgres” (now rarely in all capital letters) because of tradition or because it is easier to pronounce. This usage is widely accepted as a nickname or alias.” This thesis will use Postgres as an alias for PostgreSQL as well.

More than 30 years after the first version Postgres has been considered the most used ORDBMS for professional developers by Stack Overflow survey [4]. According to Riggs and Ciolli [2]: “The PostgreSQL feature set attracts serious users who have serious applications. Financial services companies may be PostgreSQL’s largest user group, although governments, telecommunication companies, and many other segments are strong users as well.” It is fully ACID compliant [5] and supports many kinds of data models such as relational, document, and key/value. [2]

1.2 Kubernetes

Kubernetes, also known as K8s, is an open-source platform for automating deployment, scaling, and management of containerized applications. It provides a way to manage and orchestrate containers, which are units of software that package up an application and its dependencies into a single, isolated package that can run consistently on any infrastructure. [6]

As described by Kubernetes Documentation [7] Kubernetes provides several key features, including:

- **Service discovery:** A container can be exposed by Kubernetes either through its DNS name or its own IP address.
- **Load balancing:** In the case of high traffic to a container, stability of the deployment can be ensured by Kubernetes load balancing and distributing the network traffic.
- **Storage Orchestration:** Storage orchestration in Kubernetes allows for the automatic mounting of a storage system of choice, including local storage, public cloud providers, and others.
- **Automated rollouts and rollbacks:** The desired state of deployed containers can be described using Kubernetes, and the actual state can be changed to the desired state at a controlled rate. For instance, the automation of Kubernetes can be utilized to create new containers for the deployment, remove existing containers, and transfer all their resources to the newly created container.
- **Automatic bin packing:** A cluster of nodes for running containerized tasks is provided to Kubernetes. The amount of CPU and memory required by each container is specified to Kubernetes. The optimal utilization of resources can be achieved by Kubernetes fitting the containers onto the nodes.
- **Self healing:** Containers that fail are restarted by Kubernetes, those that do not respond to the user-defined health check are replaced or killed, and they are not advertised to clients until they are deemed ready to serve.
- **Secret and configuration management:** Sensitive information, such as passwords, OAuth tokens, and SSH keys, can be stored and managed by Kubernetes. The deployment and updating of secrets and application configuration can be

done without the need to rebuild container images and without the exposure of secrets in the stack configuration.

1.2.1 Kubernetes components

Kubernetes cluster is composed of a set of worker machines that run containerized applications called nodes. Each cluster must have at least one node. [7]




Figure 1.1 The components of a Kubernetes cluster [7]

The Kubernetes control plane is the management system of a Kubernetes cluster, responsible for maintaining the desired state of the cluster. It consists of multiple components that work together to manage the cluster and its resources, including pods, services, and volumes. The key components of control plane are [8]:

- **kube-apiserver:** Acts as the front-end for the Kubernetes API and exposes the API to other components. [7]
- **etcd:** Highly available distributed key-value store that serves as the backing store for the cluster's configuration data. [9]
- **kube-scheduler:** Assigns work to nodes in the cluster, such as scheduling pods to run on nodes. [10]
- **kube-controller-manager:** Monitors the cluster's state and makes adjustments as necessary to maintain the desired state. [8]

- **cloud-controller-manager:** Manages cloud-related tasks such as node creation and management, volume management, and load balancing, allowing the other components of the control plane to focus on their specific responsibilities. Cloud manager is optional. Can be avoided when Kubernetes not used in cloud. [7]

Node components: Node components in a Kubernetes cluster run on each node and provide crucial functionality for the operation of containers on that node. [7]

- **kubelet:** Is responsible for communicating with the control plane and ensuring that containers are running and healthy. [11]
- **kube-proxy:** Is responsible for maintaining network rules on the nodes, allowing network communication to the containers. It enables the containers in a pod to communicate with other containers and the outside world, and performs tasks such as load balancing and traffic routing. [11]
- **container runtime:** Is responsible for running containers. [7]

1.2.2 Kubernetes Concepts

Pod, manifest, service replica set etc. Is something really fabulous

1.3 Running Postgres in Kubernetes

1.4 Architecture

2 NADPISY A PODNADPISY

Na této stránce je k vidění způsob tvorby různých úrovní nadpisů.

2.1 Podnadpis A

Text

2.2 Podnadpis B

Text

2.3 Podnadpis C

Text

2.3.1 Podpodnadpis alfa

Text

2.3.2 Podpodnadpis beta

Text

2.3.3 Podpodnadpis gama

Text

2.4 Podnadpis D

Text

3 VKLÁDÁNÍ OBRÁZKŮ, TABULEK A CITACÍ

Níže následují ukázky vložení obrázku, tabulky a různorodých citací.

3.1 Obrázek

Obrázek 3.1 prezentuje logo Fakulty aplikované informatiky.




Figure 3.1 Popisek obrázku

3.2 Tabulka

Tabulka 3.1 obsahuje dva řádky a celkem 7 sloupců.

Table 3.1 Popisek tabulky

	1	2	3	4	5	Cena [Kč]
F	(jedna)	(dva)	(tři)	(čtyři)	(pět)	300

II. PRAKTICKÁ ČÁST

4 NADPIS PRVNÍ KAPITOLY PRAKTICKÉ ČÁSTI

Text je text

ZÁVĚR

Text závěru.

REFERENCES

- [1] Group, T. P. G. D.: PostgreSQL 15.1 Documentation - What Is PostgreSQL? [visited 2023-02-08].
URL <https://www.postgresql.org/docs/15/intro-what-is.html>
- [2] Riggs, S.; Ciolli, G.: *PostgreSQL 14 Administration Cookbook*, chapter Introducing PostgreSQL 14. Birmingham: Packt publishing, 1st edition, 2022, ISBN 9781803248974, pp. 2–8.
- [3] Group, T. P. G. D.: PostgreSQL 15.1 Documentation - History. [visited 2023-02-08].
URL <https://www.postgresql.org/docs/15/history.html>
- [4] Inc., S. E.: Stack Overflow 2022 Developer Survey. [visited 2023-02-07].
URL <https://survey.stackoverflow.co/2022/>
- [5] Juba, S.; Vannahme, A.; Volkov, A.: *Learning PostgreSQL*, chapter Relational databases. Packt Publishing Ltd, 1st edition, 2015, ISBN 9781783989188, pp. 1–4.
- [6] Vayghan, L. A.; Saied, M. A.; Toeroe, M.; et al.: Kubernetes as an availability manager for microservice applications. *arXiv preprint arXiv:1901.04946*, 2019.
- [7] Authors, T. K.: Kubernetes Components. [visited 2023-02-07].
URL <https://kubernetes.io/docs/concepts/overview/components/>
- [8] Sayfan, G.: *Mastering Kubernetes*, chapter Kubernetes concepts. Birmingham: Packt Publishing, third edition edition, [2020]., ISBN 978-183-9211-256.
- [9] Dobies, J.; Wood, J.: *Kubernetes operators*. Sebastopol, CA: O'Reilly Media, 2020, ISBN 978-149-2048-046.
- [10] Burns, B.; Beda, J.; Hightower, K.; et al.: *Kubernetes: up and running*, chapter Pods. " O'Reilly Media, Inc.", third edition edition, 2022.
- [11] Burns, B.; Beda, J.; Hightower, K.; et al.: *Kubernetes: up and running*. " O'Reilly Media, Inc.", third edition edition, 2022.

LIST OF ABBREVIATIONS

ORDBMS	Object-relational Database Management System
ACID	Atomicity, Consistency, Isolation, Durability
K8s	Kubernetes
CPU	Central Processing Unit
PTFE	Polytetrafluoroethylene
VNA	Vector Network Analyser

LIST OF FIGURES

Fig. 1.1.	The components of a Kubernetes cluster [7]	12
Fig. 3.1.	Popisek obrázku.....	15

LIST OF TABLES

Tab. 3.1. Popisek tabulky	15
---------------------------------	----

LIST OF APPENDICES

A I. Název přílohy

APPENDIX A I. NÁZEV PŘÍLOHY

Obsah přílohy